Prune after encoding (FastV, SparseVLM, IF-Prune)

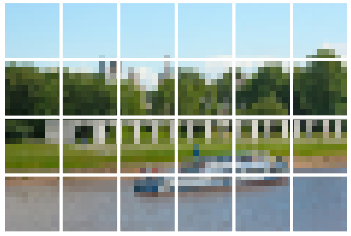

ViT: all 24 tiles → LLM: all tokens, then prune → peak memory grows with *N*

FoveaTile: look before you encode

coverage view → prior

ViT: 4 of 24 tiles → LLM: budget tokens only → peak memory constant

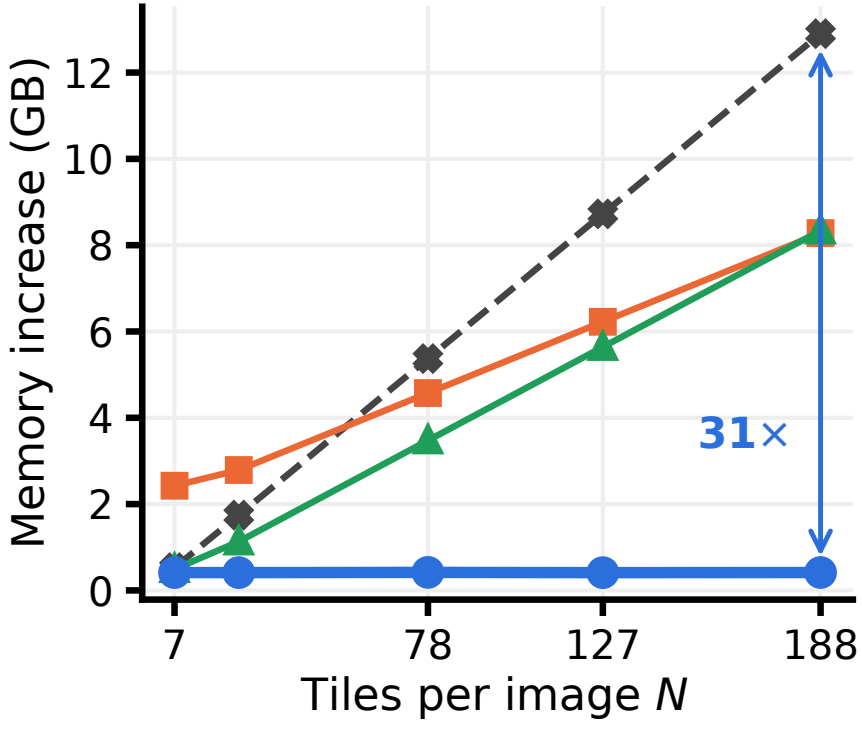

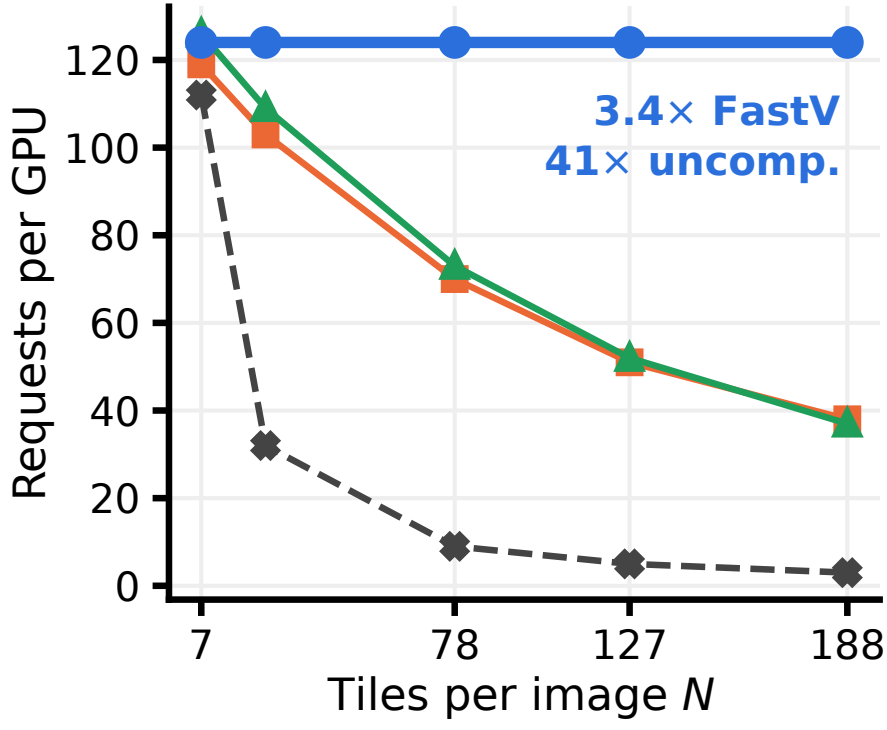

FoveaTile (ours) FastV IF-Prune Uncompressed